

Course Project

Load stuff

```
require(caret)
require(ggplot2)
require(reshape2)
require(plyr)

set.seed(42) # make results reproducible

ds <- read.csv('data/pml-training.csv')
test <- read.csv('data/pml-testing.csv')
```

Exploratory data analysis

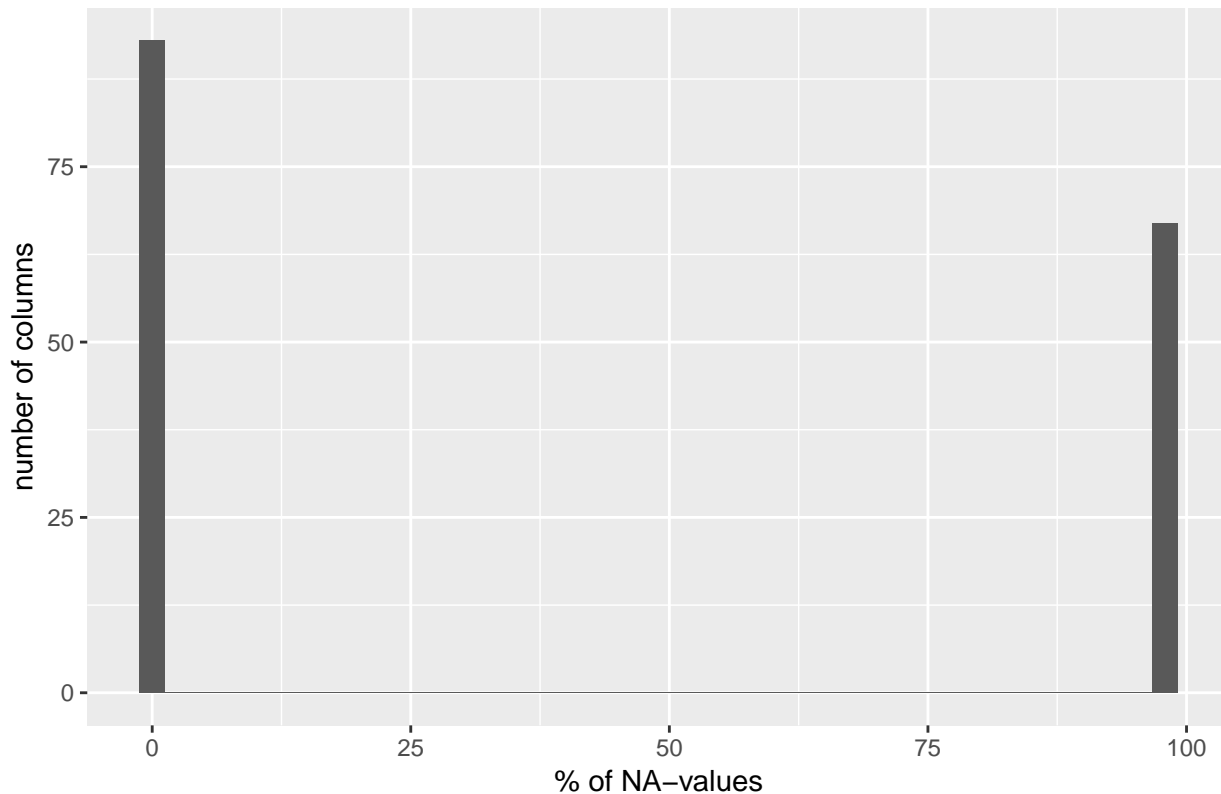
```
ds_complete <- na.omit(ds)

na_count <- apply(ds, 2, function(x) sum(is.na(x))) # count number of missing values per column
na_cols <- names(na_count[na_count != 0])
complete_cols <- names(na_count[na_count == 0])

na_per_col <- data.frame(na_count)
na_per_col$colname <- rownames(na_per_col)
na_per_col$percent <- (na_per_col$na_count / nrow(ds)) * 100

g <- ggplot(data=na_per_col) +
  geom_histogram( aes(x=percent), bins=40 ) +
  xlab('% of NA-values') +
  ylab('number of columns') +
  ggtitle('Percent of NA-values in columns') +
  theme(plot.title = element_text(size=17, face='bold', hjust=0.5))
g
```

Percent of NA-values in columns



There are 19622 observations of barbell exercises in the dataset (one observation being one repetition). The number of features is 158. However, the data is not complete. All except 406 observations contain missing values.

The missing values are not distributed uniformly over the columns. 93 columns have no missing values at all. On the other hand, 67 columns have close to 100% of missing values (almost as if they were deleted on purpose... :-). With such a high fraction of missing values, one can already suspect that 67 of the features will not be of much use for the prediction. Nevertheless, for the sake of gathering experience, I decided to analyze all features.

It has to be noted that the dataset is ordered, with all the class-A observations at the top of the dataset, followed by class-B observations and so on. Therefore, when fitting a preliminary model, the index-variable X - which just indicates the position of the observation in the dataframe - was identified as the most important 'feature'. Of course, in the test-set or a real-world dataset, the data will not be ordered by class. Therefore, I make sure not to include index variable 'X' in the feature list. For the same reason, I excluded the timestamp-columns from the features. The name of the person performing the exercise was also not used as a feature. It is possible that the name is a useful feature for prediction of the test set, because the test set was collected with the same people as the training set. However, the study aims to identify general patterns which do not depend on the person doing the exercise. Therefore, the information who performed the exercise is intentionally ignored in the model.

Functions for fitting and evaluating random forest models

```
trainRF <- function(trainingdata, features) {  
  cv_params <- trainControl(method = 'repeatedcv',  
                             number = 4, # number of folds
```

```

        repeats = 1, # number of times, the cv is repeated
        allowParallel = TRUE)
RFmodel <- train(classe ~ .,
                data = trainingdata[,c(features, 'classe')], # only features and outcome variable (no
                method = 'rf',
                trControl = cv_params,
                importance = TRUE, # for feature importance
                verbose = FALSE)
return(RFmodel)
}

get_probabilities <- function(testdata, features, my_model) {
  prob <- predict(my_model,
                  newdata = testdata[,features],
                  type = 'prob')
  prob$id <- rownames(prob)
  long <- melt(prob, id.vars='id')
  return(long[order(long$id),])
}

get_predictions <- function(probabilities) { # long format dataframe of probabilities
  pred <- ddply(probabilities, 'id', function(x) { # get prediction from probability
    return(head(subset(x, value == max(x$value)), 1)) # in case of equal probabilities: take first
  })
  colnames(pred) <- c('id', 'prediction', 'probability')
  return(pred)
}

evaluate <- function(predictions, testdata) { #
  truth <- data.frame(rownames(testdata), testdata[, 'classe'])
  colnames(truth) <- c('id', 'classe')
  compare <- merge(predictions, truth, by='id')
  accuracy <- nrow(subset(compare, prediction == classe)) / nrow(compare)
  return(list(compare, accuracy))
}

get_accuracy <- function(testdata, features, my_model) { # just calculate accuracy (predictions not ret
  prob <- get_probabilities(testdata, features, my_model)
  pred <- get_predictions(prob)
  return(evaluate(pred, testdata)[[2]])
}

```

Feature importance

A runtime of about 5 minutes on this small dataset (using 4-fold cross-validation) revealed it wouldn't be computationally feasible to use all observations and all features for fitting a model. Hence, even if the whole dataset was complete, I would have to select a subset of features.

The feature selection process below is very simple. First, a preliminary random forest model was fitted using all features and the importance of the features was determined using the `varImp`-function. Then, a model was fit to the small training set using only the top 5, top 10, top 15, etc. features.

```

#mini <- head(ds_complete, 0) # empty dataframe
#for (lev in levels(ds_complete$classe)) {
#  sub <- subset(ds_complete, classe == lev)[1:50,]
#  mini <- rbind(mini, sub) # add subset to mini-trainingset
#}
#mini <- ds_complete

# Split the small, complete dataset of 406 rows into a training set and a test set:
inTraining <- createDataPartition(ds_complete$classe, p = .75, list = FALSE)
minitrain <- ds_complete[inTraining,]
minitest <- ds_complete[-inTraining,]
#truth <- data.frame(rownames(minitest), minitest[, 'classe'])
#colnames(truth) <- c('id', 'classe')

```

good link about runtime: <https://www.quora.com/What-is-the-time-complexity-of-Random-Forest-both-building-the-model-and-predicting>

Fit preliminary model to analyze feature importance

```

RFmodel_prelim <- trainRF(minitrain, feat)

feat_imp <- data.frame(varImp(RFmodel_prelim)$importance)
feat_imp['average'] <- rowMeans(feat_imp[,c('A', 'B', 'C', 'D', 'E')])
feat_imp <- feat_imp[with(feat_imp, order(average, decreasing=TRUE)),]

top50 <- rownames(head(feat_imp, 50))
top30 <- rownames(head(feat_imp, 30))
top20 <- rownames(head(feat_imp, 20))
top10 <- rownames(head(feat_imp, 10))

head(feat_imp, 10)

```

	A	B	C	D	E
## avg_roll_dumbbell	75.86647	50.41750	95.75188	80.90973	48.21507
## stddev_roll_belt	72.15122	62.41339	64.48147	57.68397	92.26774
## var_roll_belt	68.86291	60.25790	61.31898	68.62224	83.90282
## roll_belt	77.28474	53.32106	57.95503	67.98371	82.82783
## avg_roll_belt	57.93466	57.73095	69.52614	72.72157	78.63690
## var_accel_dumbbell	58.89469	73.04282	77.76597	63.09779	62.76703
## min_roll_forearm	85.81663	45.57424	61.01663	100.00000	42.74429
## var_total_accel_belt	72.31072	56.03514	54.52745	63.03789	87.41967
## avg_yaw_belt	76.65897	68.30287	60.85587	65.03334	48.34093
## amplitude_pitch_belt	78.90988	32.19863	60.25497	61.39122	84.61509
## average					
## avg_roll_dumbbell	70.23213				
## stddev_roll_belt	69.79956				
## var_roll_belt	68.59297				
## roll_belt	67.87448				
## avg_roll_belt	67.31004				
## var_accel_dumbbell	67.11366				
## min_roll_forearm	67.03036				
## var_total_accel_belt	66.66617				
## avg_yaw_belt	63.83840				
## amplitude_pitch_belt	63.47396				

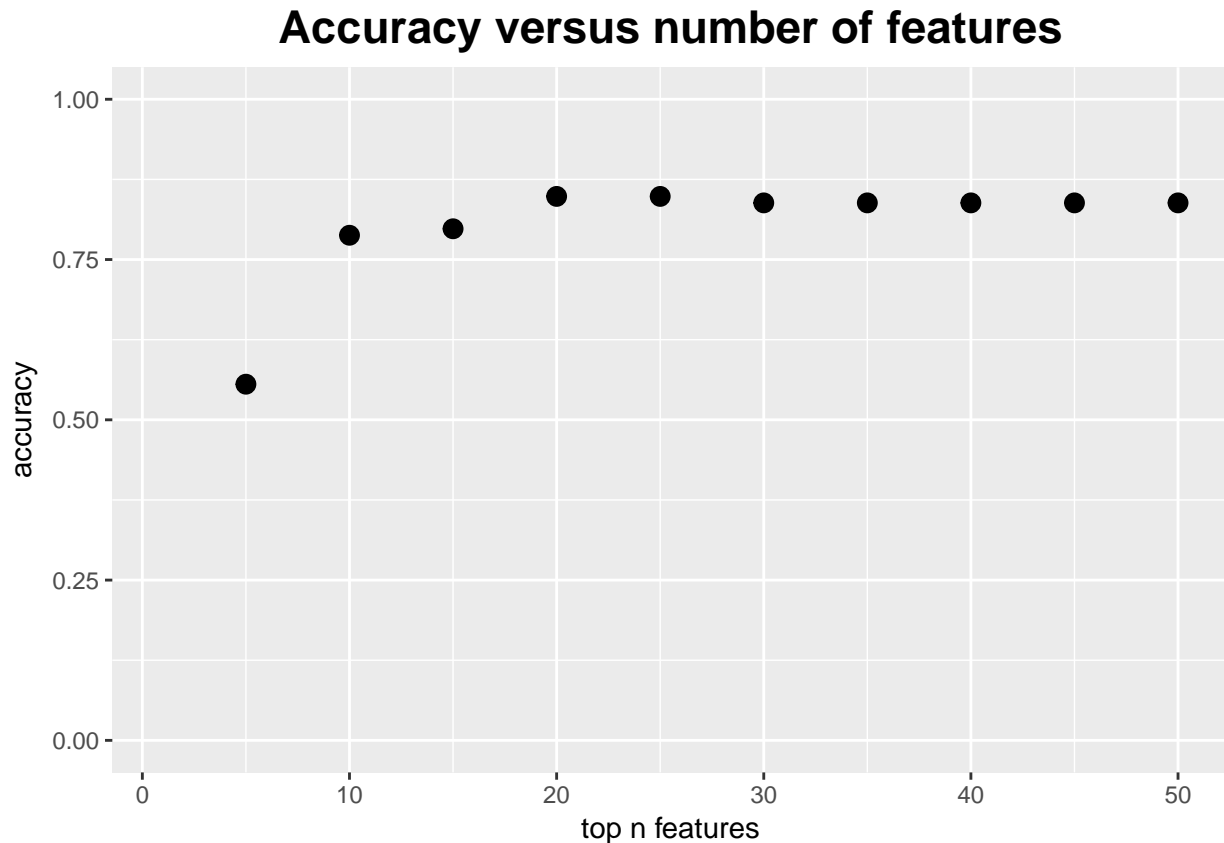
Plot accuracy as a function of feature importance

```
get_accuracy_for_feature_subsets <- function(trainingdata, testdata, feature_importance) {  
  top_n_features <- (1:10)*5 # which top n features to test  
  accuracy <- c()  
  for (i in top_n_features) {  
    top_features <- rownames(head(feature_importance, i))  
    RFmodel <- trainRF(trainingdata, top_features)  
    accuracy <- c(accuracy, get_accuracy(testdata, top_features, RFmodel))  
  }  
  return(data.frame(top_n_features, accuracy))  
}
```

```
feat_acc <- get_accuracy_for_feature_subsets(minitrain, minitest, feat_imp)  
feat_acc
```

```
##   top_n_features  accuracy  
## 1              5 0.5555556  
## 2             10 0.7878788  
## 3             15 0.7979798  
## 4             20 0.8484848  
## 5             25 0.8484848  
## 6             30 0.8383838  
## 7             35 0.8383838  
## 8             40 0.8383838  
## 9             45 0.8383838  
## 10            50 0.8383838
```

```
g <- ggplot(data=feat_acc, aes(x=top_n_features, y=accuracy)) +  
  geom_point(size=3) +  
  scale_y_continuous(limits = c(0, 1)) +  
  scale_x_continuous(limits = c(1, 50)) +  
  xlab('top n features') +  
  ggtitle('Accuracy versus number of features') +  
  theme(plot.title = element_text(size=17, face='bold', hjust=0.5))  
g
```



The plot above shows, that the accuracy increases with the use of more features until about 20 features. Then, the addition of more features does not increase the accuracy. The small differences in accuracy when using 20 or more features represent only the random fluctuations during fitting of a random forest model and can be safely ignored.

k-nearest neighbour imputation

Due large fraction of missing values in many features, imputation cannot be expected to work very well. Nevertheless, a k-nearest neighbour imputation model was fitted on the small dataset of complete values (406 rows). Then, using this model, the missing value in the whole dataset (19622 rows) were imputed.

```
knn_imp_model <- preProcess(ds_complete, method = 'knnImpute')
ds_imp <- predict(knn_imp_model, ds)
```

Split whole dataset into training and validation set

```
inTraining <- createDataPartition(ds$classe, p = .75, list = FALSE) # split the training data-set in
```

Fit model with imputed features

```
top_imputed_features <- top20
training_imp <- ds_imp[inTraining,]
```

```

validation_imp <- ds_imp[-inTraining,]

RFmodel_imp <- trainRF(training_imp, top_imputed_features)
prob <- get_probabilities(validation_imp, top_imputed_features, RFmodel_imp)
pred <- get_predictions(prob)
eval <- evaluate(pred, validation_imp)
acc <- eval[[2]]
acc

## [1] 0.9692088

```

Fit model with real values only

```

feat_complete <- feat[feat %in% complete_cols] # features without missing values
top_complete_features <- top50[top50 %in% complete_cols]

training <- ds[inTraining,]
validation <- ds[-inTraining,]

RFmodel <- trainRF(training, top_complete_features)
prob <- get_probabilities(validation, top_complete_features, RFmodel)
pred <- get_predictions(prob)
eval <- evaluate(pred, validation)
acc <- eval[[2]]
acc

## [1] 0.9781811

```